

Redis Deployment Patterns

Minimal Redis Deployment Patterns

Context: For use in early-stage, agent-driven DevOps workflows and LLM orchestration requiring lightweight Redis deployment on Kubernetes (using kind cluster).

Purpose of This Guide

This guide provides two minimal and valid Redis deployment examples for Kubernetes that can:

- Run on modest local machines (2 CPUs, 8GB RAM)
- Be used for lightweight caching, session memory, or simple agent memory
- Act as stable grounding data for LLMs building infrastructure plans
- The use cases are valid for low to high throughput workloads with 80% reads and 20% writes access pattern.

Configurations

Setting	Value
Redis Image	redis:7.2
Deployment Type	apps/v1 Deployment
Storage	Ephemeral only
Persistence	Disabled
Platform	Kubernetes
CPU/Memory Budget	2 CPUs / 2Gi RAM
Network	Exposed via NodePort

Deployment Example 1: Single-Replica Redis for Low throughput

For basic local caching and memory store use cases. (Throughput (100 ops / second))

apiVersion: apps/v1

kind: Deployment

metadata:

 name: redis-deployment

 labels:

 app: redis

spec:

 replicas: 1

 selector:

 matchLabels:

 app: redis

 template:

 metadata:

 labels:

 app: redis

 spec:

 containers:

 - name: redis

 image: redis:7.2

 ports:

 - containerPort: 6379

 resources:

 requests:

memory: "512Mi"

cpu: "250m"

limits:

memory: "1Gi"

cpu: "500m"

apiVersion: v1

kind: Service

metadata:

name: redis-service

spec:

type: NodePort

selector:

app: redis

ports:

- protocol: TCP

port: 6379

targetPort: 6379

nodePort: 30079

Deployment Example 2: High Throughput Redis (vertical Scaling for high throughput 1000 ops/second)

For read-heavy use cases or early agentic system simulation with redundancy.

apiVersion: apps/v1

kind: Deployment

metadata:

name: redis-deployment

labels:

app: redis

spec:

replicas: 2

selector:

matchLabels:

app: redis

template:

metadata:

labels:

app: redis

spec:

containers:

- name: redis

image: redis:7.2

ports:

- containerPort: 6379

resources:

requests:

memory: "768Mi"

cpu: "500m"

limits:

memory: "1.5Gi"

cpu: "1"

apiVersion: v1

kind: Service

metadata:

name: redis-service

spec:

type: NodePort

selector:

app: redis

ports:

- protocol: TCP

port: 6379

targetPort: 6379

nodePort: 30079